

# Business Analytics - ETC3250 2017 - Lab 5 solutions

*Souhaib Ben Taieb*

*23 August 2017*

## Purpose

This lab will be on classification methods, including logistic regression, linear and quadratic discriminant analysis and nearest neighbors.

Note: the online version of “An Introduction to Statistical Learning with Applications in R” (ISLR) is available at <http://www-bcf.usc.edu/~gareth/ISL/>.

## Exercise 1

Read and run the code in Sections 4.6.1 to 4.6.6 of ISLR.

## Assignment - Question 1

Do the exercise 7 in chapter 4.7 of ISLR.

Let  $p_k(x)$  be the probability that a company will ( $k = 1$ ) or will not ( $k = 0$ ) issue a dividend this year given that its percentage profit was  $x$  last year.

Using Bayes theorem and since we assume  $X$  follows a normal distribution, we can write:

$$p_k(x) = \frac{\pi_k \frac{1}{\sqrt{2\pi}\sigma} \exp(-\frac{1}{2\sigma^2}(x - \mu_k)^2)}{\sum_{l=1}^k \pi_l \frac{1}{\sqrt{2\pi}\sigma} \exp(-\frac{1}{2\sigma^2}(x - \mu_l)^2)}, \quad k = 1, 2.$$

Then, using  $\pi_1 = .8$ ,  $\sigma = 6$ ,  $\mu_1 = 10$  and  $\mu_2 = 0$ , we have

$$p_1(x) = \frac{0.80 \exp(-\frac{1}{2*36}(x - 10)^2)}{0.80 \exp(-\frac{1}{2*36}(x - 10)^2) + 0.20 \exp(-\frac{1}{2*36}x^2)}.$$

Finally, since  $x = 4$ , we have

$$p_1(4) \approx 75\%.$$

## Assignment - Question 2

Do the exercise 8 in chapter 4.7 of ISLR.

$E_{\text{train}} = 0.20$ ,  $E_{\text{test}} = 0.30$  and  $E_{\text{avg}} = 0.25$ .


$E_{\text{train}} = x_1$ ,  $E_{\text{test}} = x_2$  and  $E_{\text{avg}} = 0.18$ .

I would prefer logistic regression since the test error rate is not too far from the training error rate. For 1-nearest neighbors, which is typically a high variance classifier, we could obtain an average error rate of 0.18 with  $x_1 = 0$  and  $x_2 = 0.36$ , which is an overfitting classifier with test error rate larger than 0.30.

### Assignment - Question 3

Do the exercise 10 in chapter 4.7 of ISL

```
library(ISLR)
summary(Weekly)
#      Year      Lag1      Lag2      Lag3
# Min.   :1990   Min.   :-18.1950   Min.   :-18.1950   Min.   :-18.1950
# 1st Qu.:1995   1st Qu.: -1.1540   1st Qu.: -1.1540   1st Qu.: -1.1580
# Median :2000   Median :  0.2410   Median :  0.2410   Median :  0.2410
# Mean   :2000   Mean   :  0.1506   Mean   :  0.1511   Mean   :  0.1472
# 3rd Qu.:2005   3rd Qu.:  1.4050   3rd Qu.:  1.4090   3rd Qu.:  1.4090
# Max.   :2010   Max.   : 12.0260   Max.   : 12.0260   Max.   : 12.0260
#      Lag4      Lag5      Volume
# Min.   :-18.1950   Min.   :-18.1950   Min.   :0.08747
# 1st Qu.: -1.1580   1st Qu.: -1.1660   1st Qu.:0.33202
# Median :  0.2380   Median :  0.2340   Median :1.00268
# Mean   :  0.1458   Mean   :  0.1399   Mean   :1.57462
# 3rd Qu.:  1.4090   3rd Qu.:  1.4050   3rd Qu.:2.05373
# Max.   : 12.0260   Max.   : 12.0260   Max.   :9.32821
#      Today      Direction
# Min.   :-18.1950   Down:484
# 1st Qu.: -1.1540   Up  :605
# Median :  0.2410
# Mean   :  0.1499
# 3rd Qu.:  1.4050
# Max.   : 12.0260
pairs(Weekly)
```



```
attach(Weekly)
```

```
glm.fit <- glm(Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 + Volume, data = Weekly, family = binomial)
summary(glm.fit)
```

```

#
# Call:
# glm(formula = Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 +
#       Volume, family = binomial, data = Weekly)
#
# Deviance Residuals:
#      Min       1Q   Median       3Q      Max
# -1.6949  -1.2565   0.9913   1.0849   1.4579
#
# Coefficients:
#              Estimate Std. Error z value Pr(>|z|)
# (Intercept)  0.26686    0.08593   3.106  0.0019 **
# Lag1        -0.04127    0.02641  -1.563  0.1181
# Lag2         0.05844    0.02686   2.175  0.0296 *
# Lag3        -0.01606    0.02666  -0.602  0.5469
# Lag4        -0.02779    0.02646  -1.050  0.2937
# Lag5        -0.01447    0.02638  -0.549  0.5833
# Volume      -0.02274    0.03690  -0.616  0.5377
# ---
# Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
#
# (Dispersion parameter for binomial family taken to be 1)
#
#      Null deviance: 1496.2  on 1088  degrees of freedom
# Residual deviance: 1486.4  on 1082  degrees of freedom
# AIC: 1500.4
#
# Number of Fisher Scoring iterations: 4

```

```

glm.probs <- predict(glm.fit, type = "response")
glm.pred <- rep("Down", length(glm.probs))
glm.pred[glm.probs > 0.5] <- "Up"
table(glm.pred, Weekly$Direction)
#
# glm.pred Down Up
#      Down   54 48
#      Up    430 557

```

```

train <- (Year <= 2008)
Weekly.test <- Weekly[!train, ]
glm.fit <- glm(Direction ~ Lag2, data = Weekly, family = binomial, subset = train)
glm.probs <- predict(glm.fit, Weekly.test, type = "response")
glm.pred <- rep("Down", length(glm.probs))
glm.pred[glm.probs > 0.5] <- "Up"
Direction.test <- Direction[!train]
table(glm.pred, Direction.test)
#      Direction.test
# glm.pred Down Up
#      Down     9  5
#      Up     34 56
mean(glm.pred == Direction.test)
# [1] 0.625

```

```
library(MASS)
lda.fit = lda(Direction ~ Lag2, data = Weekly, subset = train)
lda.pred = predict(lda.fit, Weekly.test)
table(lda.pred$class, Direction.test)
#      Direction.test
#      Down Up
# Down    9  5
# Up     34 56
mean(lda.pred$class == Direction.test)
# [1] 0.625
```

```
qda.fit = qda(Direction ~ Lag2, data = Weekly, subset = train)
qda.class = predict(qda.fit, Weekly.test)$class
table(qda.class, Direction.test)
#      Direction.test
# qda.class Down Up
# Down    0  0
# Up     43 61
mean(qda.class == Direction.test)
# [1] 0.5865385
```

```
library(class)
train.X = as.matrix(Lag2[train])
test.X = as.matrix(Lag2[!train])
train.Direction = Direction[train]
set.seed(1)
knn.pred = knn(train.X, test.X, train.Direction, k = 100, prob = T)
table(knn.pred, Direction.test)
#      Direction.test
# knn.pred Down Up
# Down    10 11
# Up     33 50
```

```
glm.fit = glm(Direction ~ Lag2:Lag1, data = Weekly, family = binomial, subset = train)
glm.probs = predict(glm.fit, Weekly.test, type = "response")
glm.pred = rep("Down", length(glm.probs))
glm.pred[glm.probs > 0.5] = "Up"
Direction.0910 = Direction[!train]
table(glm.pred, Direction.test)
#      Direction.test
# glm.pred Down Up
# Down    1  1
# Up     42 60
mean(glm.pred == Direction.test)
# [1] 0.5865385
```

```
lda.fit = lda(Direction ~ Lag2:Lag1, data = Weekly, subset = train)
lda.pred = predict(lda.fit, Weekly.test)
mean(lda.pred$class == Direction.test)
# [1] 0.5769231
```

```

for(k in seq(1, 100, by = 10)){
  knn.pred = knn(train.X, test.X, train.Direction, k = k, prob = T)
  print(table(knn.pred, Direction.test))
}
#           Direction.test
# knn.pred Down Up
#   Down    21 30
#   Up      22 31
#           Direction.test
# knn.pred Down Up
#   Down    18 21
#   Up      25 40
#           Direction.test
# knn.pred Down Up
#   Down    19 21
#   Up      24 40
#           Direction.test
# knn.pred Down Up
#   Down    19 25
#   Up      24 36
#           Direction.test
# knn.pred Down Up
#   Down    21 24
#   Up      22 37
#           Direction.test
# knn.pred Down Up
#   Down    20 23
#   Up      23 38
#           Direction.test
# knn.pred Down Up
#   Down    20 18
#   Up      23 43
#           Direction.test
# knn.pred Down Up
#   Down    12 15
#   Up      31 46
#           Direction.test
# knn.pred Down Up
#   Down     9 10
#   Up      34 51
#           Direction.test
# knn.pred Down Up
#   Down    10 10
#   Up      33 51

```

## TURN IN

- Your .Rmd file (which should knit without errors and without assuming any packages have been pre-loaded)
- Your Word (or pdf) file that results from knitting the Rmd.
- DUE: August 27, 11:55pm (late submissions not allowed), loaded into moodle